Щербань Георгий
Разработка и развертывание ИИ-приложения в Kubernetes

1. Введение

Цель проекта — разработать и развернуть в локальном Kubernetes-кластере (Minikube) контейнеризированное Java-приложение с интегрированной ИИ-логикой анализа тональности текста.

В рамках работы были выполнены следующие задачи:

разработка REST API на Spring Boot

контейнеризация с помощью Docker

развертывание в Minikube (Deployment + Service + Ingress)

настройка мониторинга (Prometheus + Actuator)

настройка горизонтального автоскейлинга HPA

проведение нагрузочного тестирования

анализ трендов в области ИИ, контейнеризации и оркестрации

Используемый стек:

1. Spring Boot 3.3, Java 17
2. Docker, multi-stage build
3. Minikube (драйвер Docker)
4. kubectl, helm, metrics-server
5. Horizontal Pod Autoscaler (HPA)
6. Nginx Ingress Controller

2. Развёртывание инфраструктуры Minikube

2.1 Установка и запуск кластера

Кластер Minikube был развёрнут на Windows 11 с драйвером Docker.

minikube start --driver=docker --cpus=3 --memory=3500mb --nodes=1

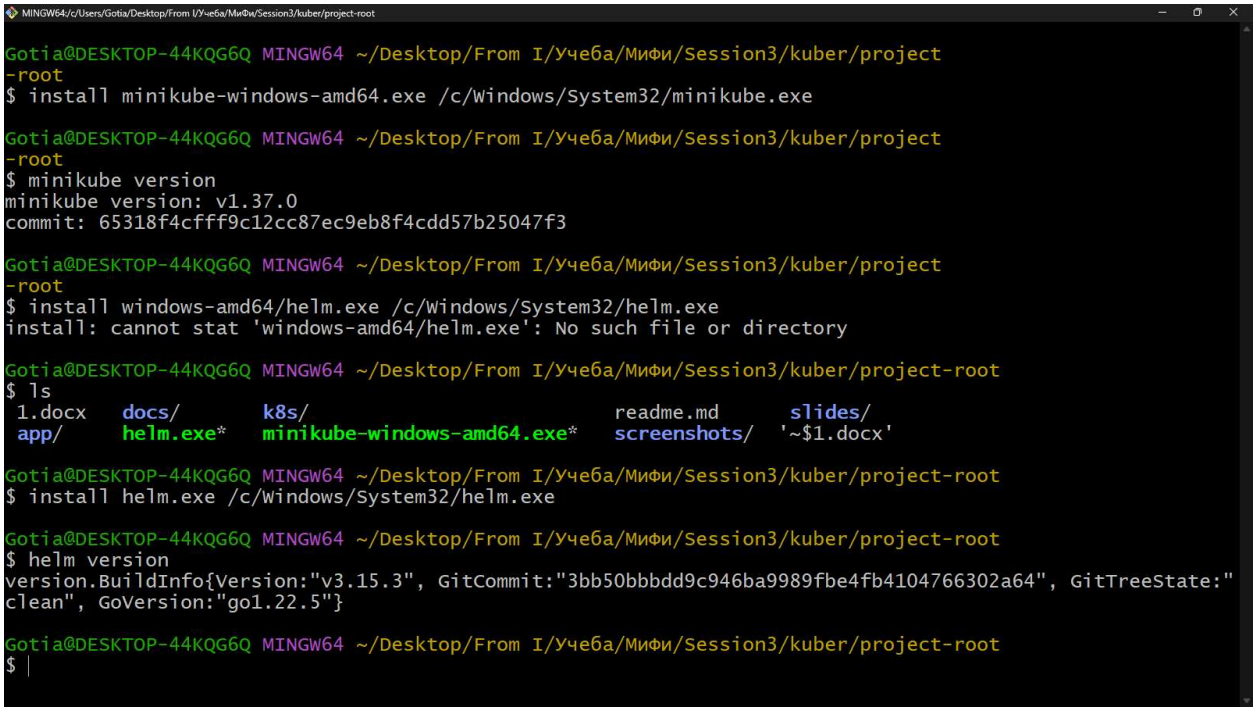

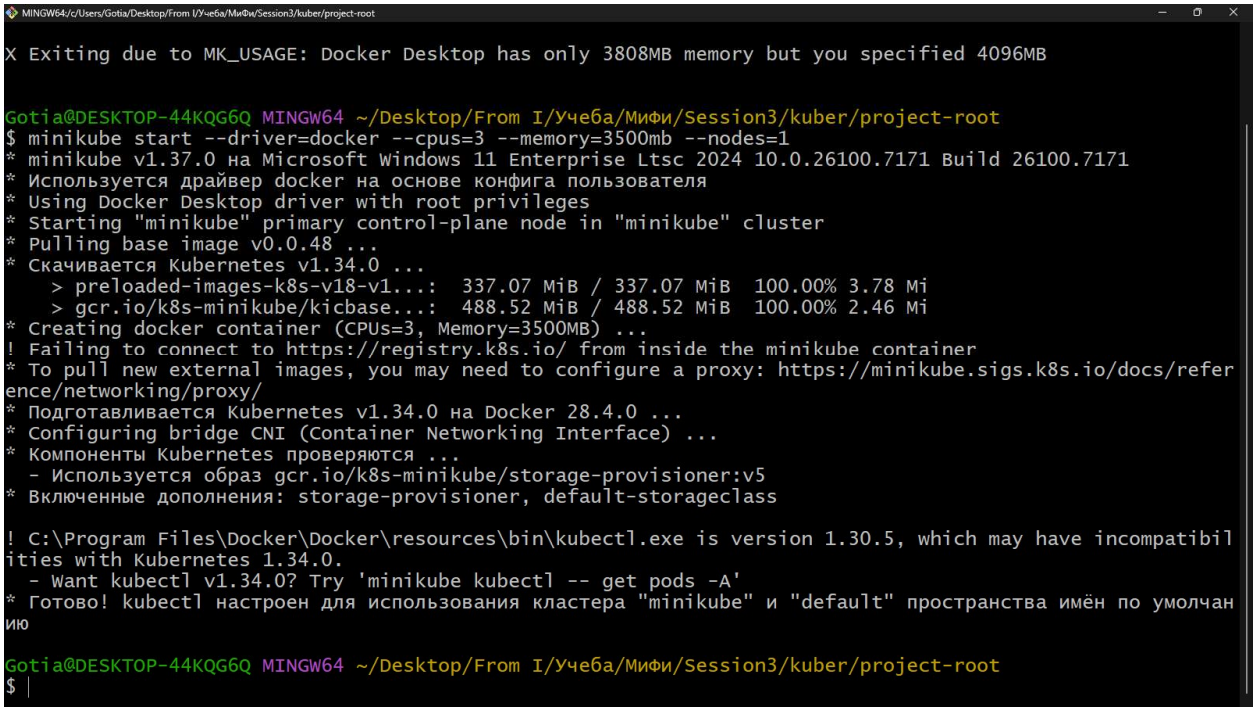

Ограничение памяти обусловлено настройками Docker Desktop, но конфигурации достаточно для HPA, ingress и сервиса.

2.2 Включение необходимых аддонов

minikube addons enable ingress

minikube addons enable metrics-server

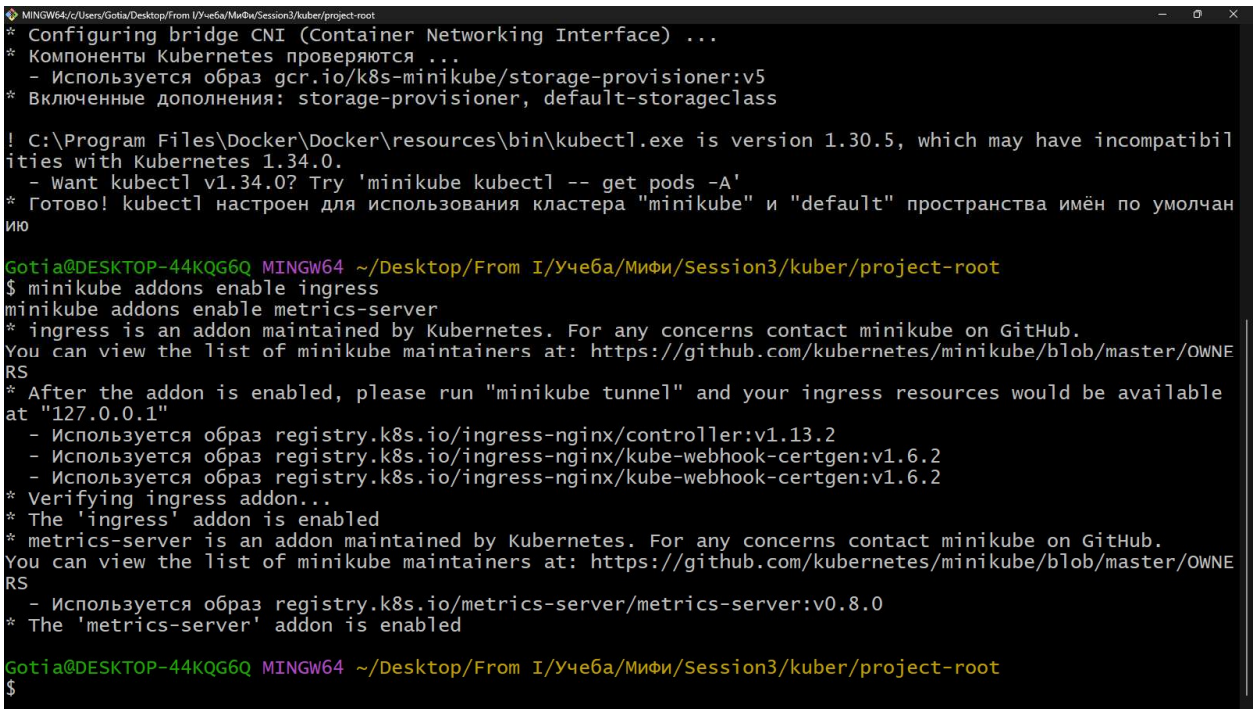

Ingress обеспечивает маршрутизацию запросов,
metrics-server - обязательное условие для работы HPA.

2.3 Проверка состояния кластера

kubectl get pods -A

kubectl get nodes

minikube status

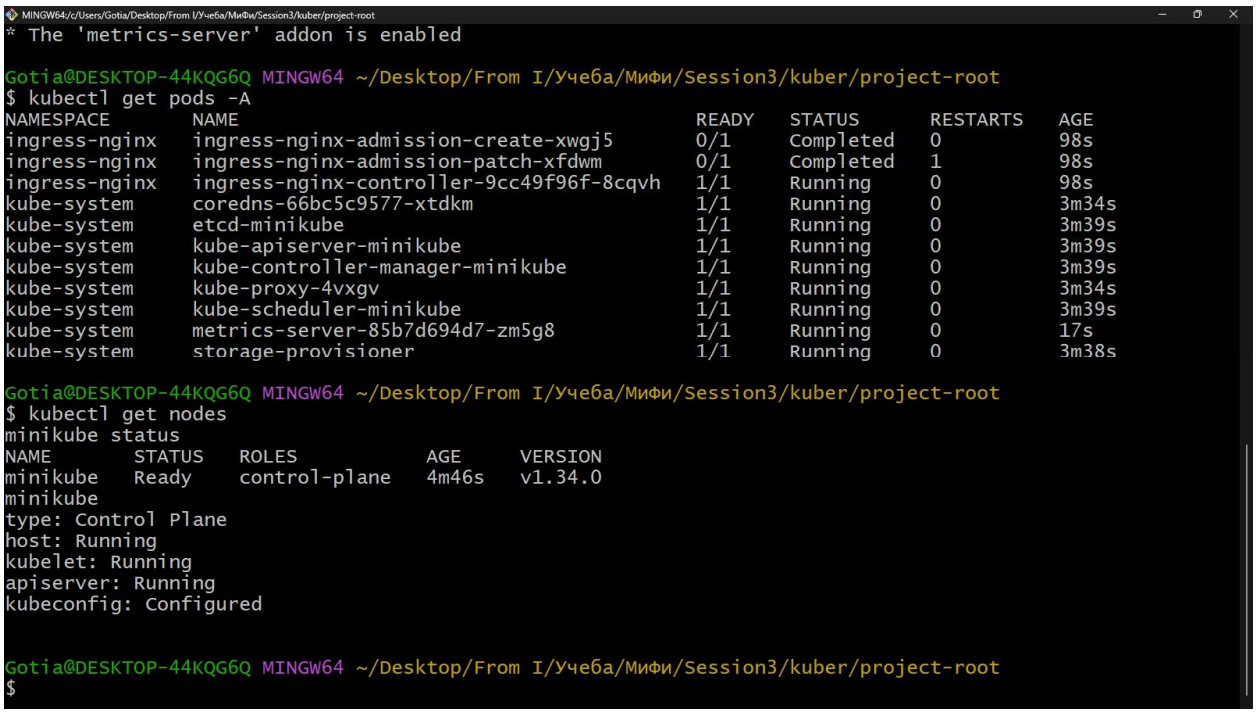

ingress-nginx-controller работает,

metrics-server функционирует,

нода находится в состоянии Ready.

3. Разработка Java-приложения

Spring Boot приложение с REST-эндпоинтом:

GET /api/sentiment?text=...

Пример ответа:

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ curl "http://localhost:8080/api/sentiment?text=I%20love%20kubernetes"
{"sentiment":"positive","input":"I love kubernetes"}
```

запуск:

mvn clean package -DskipTests

mvn spring-boot:run

Проверка:

curl http://localhost:8080/api/sentiment?text=I%20love%20kubernetes

```
[INFO] Copying 1 resource from src\main\resources to target\classes
[INFO]
[INFO] --- compiler:3.13.0:compile (default-compile) @ sentiment-app ---
[INFO] Recompiling the module because of changed source code.
[WARNING] File encoding has not been set, using platform encoding windows-1251, i.e. build is platform dependent!
[INFO] Compiling 2 source files with javac [debug target 1.8] to target\classes
[WARNING] bootstrap class path not set in conjunction with -source 8
[INFO]
[INFO] --- resources:3.3.1:testResources (default-testResources) @ sentiment-app ---
[WARNING] Using platform encoding (Cp1251 actually) to copy filtered resources, i.e. build is platform dependent!
[INFO] skip non existing resourceDirectory C:\Users\Gotia\Desktop\From I\Учеба\МиФи\Session3\kuber\project-root\app\SentimentApplication\src\test\resources
[INFO]
[INFO] --- compiler:3.13.0:testCompile (default-testCompile) @ sentiment-app ---
[INFO] Recompiling the module because of changed dependency.
[WARNING] File encoding has not been set, using platform encoding windows-1251, i.e. build is platform dependent!
[INFO] Compiling 1 source file with javac [debug target 1.8] to target\test-classes
[WARNING] bootstrap class path not set in conjunction with -source 8
[INFO]
[INFO] --- surefire:3.2.5:test (default-test) @ sentiment-app ---
[INFO] Tests are skipped.
[INFO]
[INFO] --- jar:3.4.1:jar (default-jar) @ sentiment-app ---
[INFO] Building jar: C:\Users\Gotia\Desktop\From I\Учеба\МиФи\Session3\kuber\project-root\app\SentimentApplication\target\sentiment-app-1.0.0.jar
[INFO] ------------------------------------------------------------------------
[INFO] BUILD SUCCESS
[INFO] ------------------------------------------------------------------------
[INFO] Total time:  15.981 s
[INFO] Finished at: 2025-11-24T21:44:59+03:00
[INFO] ------------------------------------------------------------------------

Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/From I/Учеба/МиФи/Session3/kuber/project-root/app/SentimentApplication (master)
```

```
[INFO] --- compiler:3.13.0:testCompile (default-testCompile) @ sentiment-app ---
[INFO] Nothing to compile - all classes are up to date.
[INFO]
[INFO] <<< spring-boot:4.0.0:run (default-cli) < test-compile @ sentiment-app <<<
[INFO]
[INFO]
[INFO] --- spring-boot:4.0.0:run (default-cli) @ sentiment-app ---
[INFO] Attaching agents: []

  .   ____          _            __ _ _
 /\\ / ___'_ __ _ _(_)_ __  __ _ \ \ \ \
( ( )\___ | '_ | '_| | '_ \/ _` | \ \ \ \
 \\/  ___)| |_)| | | | | || (_| |  ) ) ) )
  '  |____| .__|_| |_|_| |_\__, | / / / /
 =========|_|==============|___/=/_/_/_/

 :: Spring Boot ::                (v3.3.4)

2025-11-24T21:46:45.085+03:00  INFO 5460 --- [           main] c.k.m.s.SentimentApplication             : Starting SentimentApplication using Java 17.0.12 with PID 5460 (C:\Users\Gotia\De
sktop\From I\╚ÿxcp\╠╗╚╗\Session3\kuber\project-root\app\SentimentApplication\target\classes started by Gotia in C:\Users\Gotia\Desktop\From I\╚ÿxcp\╠╗╚╗\Session3\kuber\project-root\app\Se
ntimentApplication)
2025-11-24T21:46:45.104+03:00  INFO 5460 --- [           main] c.k.m.s.SentimentApplication             : No active profile set, falling back to 1 default profile: "default"
2025-11-24T21:46:46.942+03:00  INFO 5460 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer  : Tomcat initialized with port 8080 (http)
2025-11-24T21:46:46.963+03:00  INFO 5460 --- [           main] o.apache.catalina.core.StandardService   : Starting service [Tomcat]
2025-11-24T21:46:46.963+03:00  INFO 5460 --- [           main] o.apache.catalina.core.StandardEngine    : Starting Servlet engine: [Apache Tomcat/10.1.30]
2025-11-24T21:46:47.035+03:00  INFO 5460 --- [           main] o.a.c.c.C.[Tomcat].[localhost].[/]       : Initializing Spring embedded WebApplicationContext
2025-11-24T21:46:47.036+03:00  INFO 5460 --- [           main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 1780 ms
2025-11-24T21:46:47.524+03:00  INFO 5460 --- [           main] o.s.b.a.e.web.EndpointLinksResolver      : Exposing 2 endpoints beneath base path '/actuator'
2025-11-24T21:46:47.591+03:00  INFO 5460 --- [           main] o.s.b.w.embedded.tomcat.TomcatWebServer  : Tomcat started on port 8080 (http) with context path '/'
2025-11-24T21:46:47.606+03:00  INFO 5460 --- [           main] c.k.m.s.SentimentApplication             : Started SentimentApplication in 3.497 seconds (process running for 4.124)
```

4. Контейнеризация приложения

Использовал multi-stage Dockerfile:
dockerfile можно увидеть в проекте

сборка:

docker build -t sentiment-app:1.0.0 .

Проверка:

docker run --rm -p 8081:8080 sentiment-app:1.0.0

Проверка API и метрик:

curl http://localhost:8081/api/sentiment?text=hello

curl http://localhost:8081/actuator/prometheus

5. Работа с Docker Minikube и загрузка образа

minikube image load sentiment-app:1.0.0

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ minikube status
minikube
type: Control Plane
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured
```

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ minikube docker-env
You can further specify your shell with either 'cmd' or 'powershell' with the --shell flag.

SET DOCKER_TLS_VERIFY=1
SET DOCKER_HOST=tcp://127.0.0.1:58223
SET DOCKER_CERT_PATH=C:\Users\Gotia\.minikube\certs
SET MINIKUBE_ACTIVE_DOCKERD=minikube
REM To point your shell to minikube's docker-daemon, run:
REM @FOR /f "tokens=*" %i IN ('minikube -p minikube docker-env --shell cmd') DO @%i

Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ eval $(minikube -p minikube docker-env)
bash: unexpected EOF while looking for matching `''

Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ eval $(minikube -p minikube docker-env --shell=bash)

Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ echo $DOCKER_HOST
tcp://127.0.0.1:58223
```

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ minikube image load sentiment-app:1.0.0
```

6. Kubernetes-манифесты

Все ямл-файлы были подготовлены вручную.

6.1 Deployment (3 реплики)

deployment.yaml

6.2 Service LoadBalancer

service.yaml

6.3 Ingress

ingress.yaml

6.4 HPA

hpa.yaml

7. Применение манифестов

kubectl apply -f deployment.yaml

kubectl apply -f service.yaml

kubectl apply -f ingress.yaml

kubectl apply -f hpa.yaml

```
deployment.apps/sentiment-deployment created
service/sentiment-service created
ingress.networking.k8s.io/sentiment-ingress created
horizontalpodautoscaler.autoscaling/sentiment-hpa created
```

Проверка подов:

kubectl get pods

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root/k8s
$ kubectl get pods
NAME                                   READY   STATUS    RESTARTS   AGE
sentiment-deployment-7fb479685f-mrnss   1/1     Running   0          32s
sentiment-deployment-7fb479685f-w96zm   1/1     Running   0          32s
sentiment-deployment-7fb479685f-xsphc   1/1     Running   0          32s
```

Проверка сервисов:

kubectl get svc

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root/k8s
$ kubectl get svc
NAME                TYPE           CLUSTER-IP      EXTERNAL-IP   PORT(S)        AGE
kubernetes          ClusterIP      10.96.0.1       <none>        443/TCP        93m
sentiment-service   LoadBalancer   10.108.94.229   <pending>     80:32110/TCP   69s
```

8. Проверка ingress (через port-forward)

Для Windows Ingress через port 80 требует туннеля или port-forward.

kubectl port-forward svc/ingress-nginx-controller -n ingress-nginx 8080:80

curl http://localhost:8080/api/sentiment?text=hello

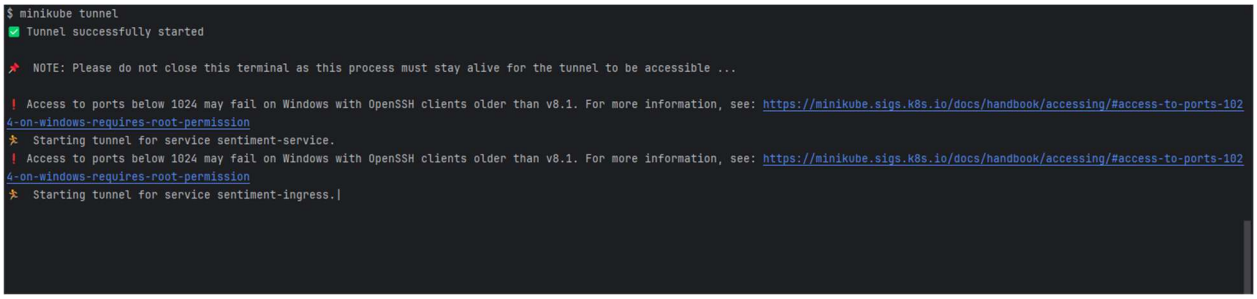

Проверка тунеля

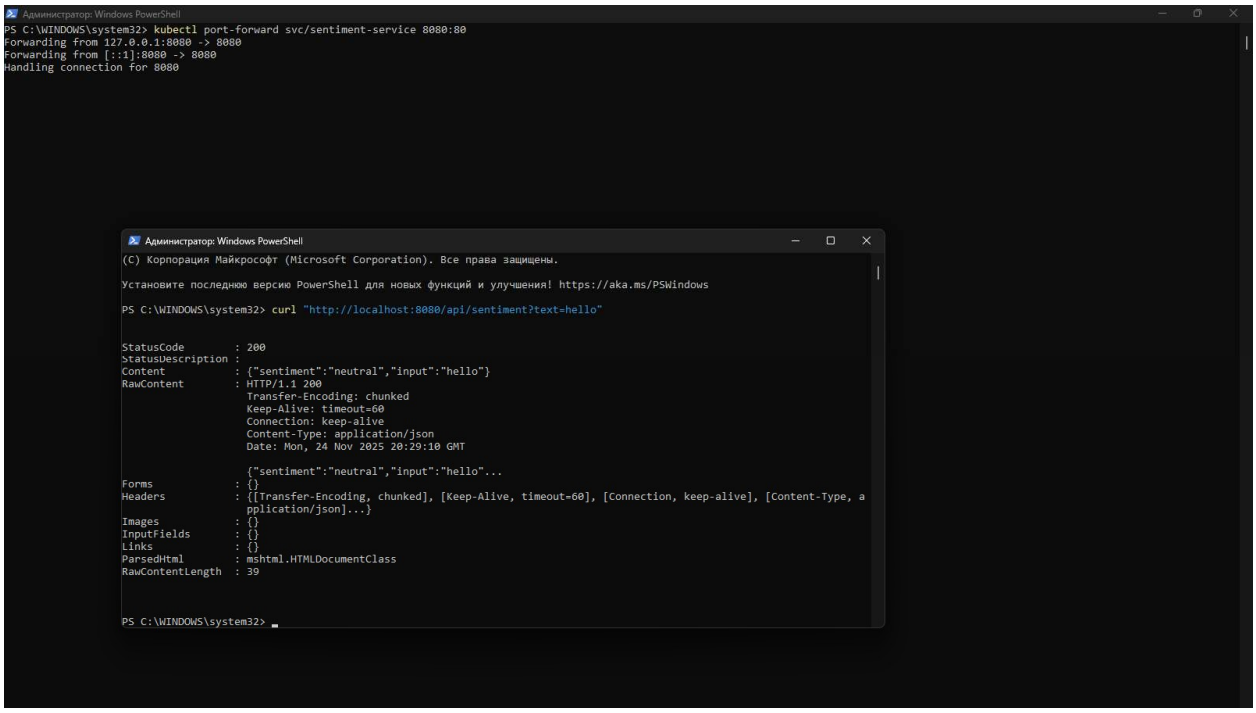

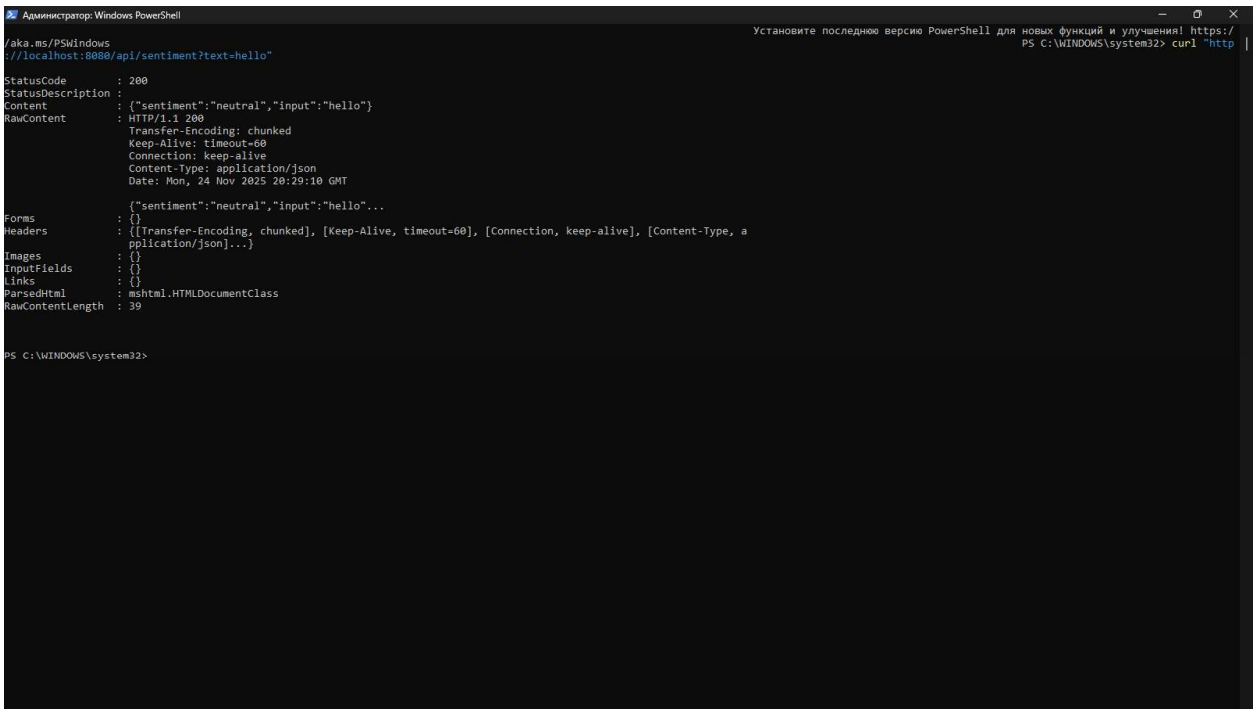
Администратор: Windows PowerShell
Установите последнюю версию PowerShell для новых функций и улучшения! https:/
/aka.ms/PSWindows
PS C:\WINDOWS\system32> curl "http
://localhost:8080/api/sentiment?text=hello"
StatusCode : 200
StatusDescription :
Content : {"sentiment":"neutral","input":"hello"}
RawContent : HTTP/1.1 200
Transfer-Encoding: chunked
Keep-Alive: timeout=60
Connection: keep-alive
Content-Type: application/json
Date: Mon, 24 Nov 2025 20:29:10 GMT
{"sentiment":"neutral","input":"hello"...
Forms : {}
Headers : {[Transfer-Encoding, chunked], [Keep-Alive, timeout=60], [Connection, keep-alive], [Content-Type, a
pplication/json]...}
Images : {}
InputFields : {}
Links : {}
ParsedHtml : mshtml.HTMLDocumentClass
RawContentLength : 39
PS C:\WINDOWS\system32>

9. HPA и нагрузочное тестирование

Подготовка контейнера:

kubectl run -it --rm load-generator --image=alpine --restart=Never -- /bin/sh

apk add --no-cache wget

Запуск бесконечной нагрузки:

while true; do wget -q -O- http://sentiment-service/api/sentiment?text=hello > /dev/null; done

Проверка:

kubectl get hpa

Результат:

TARGETS: cpu: 194%/50%

REPLICAS: 6

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ kubectl get hpa
NAME             REFERENCE                           TARGETS         MINPODS   MAXPODS   REPLICAS   AGE
sentiment-hpa    Deployment/sentiment-deployment     cpu: 194%/50%   3         6         6          29m
```

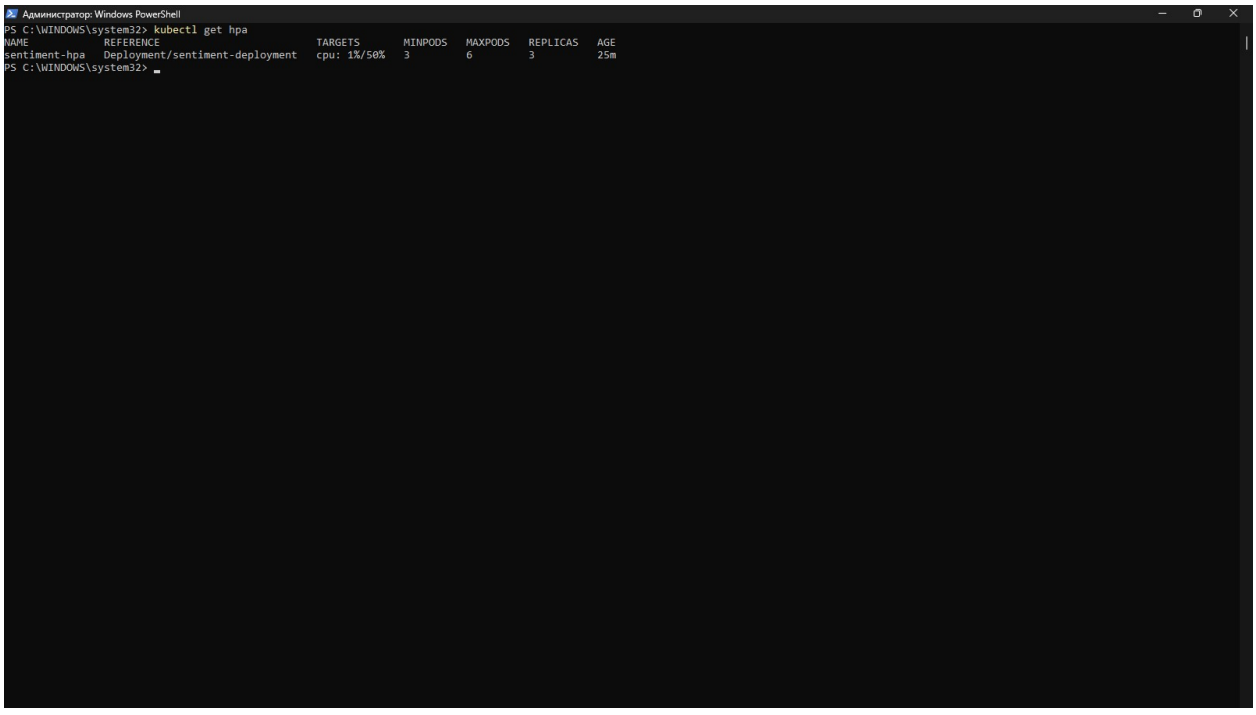

kubectl get pods

Поды увеличились с 3 до 6.

```
Gotia@DESKTOP-44KQG6Q MINGW64 ~/Desktop/CCODE/project-root
$ kubectl get pods
NAME                                   READY   STATUS    RESTARTS   AGE
load-generator                         1/1     Running   0          2m11s
sentiment-deployment-7fb479685f-6dmfc  1/1     Running   0          66s
sentiment-deployment-7fb479685f-hc8b7  1/1     Running   0          66s
sentiment-deployment-7fb479685f-mrnss  1/1     Running   0          30m
sentiment-deployment-7fb479685f-w96zm  1/1     Running   0          30m
sentiment-deployment-7fb479685f-xsphc  1/1     Running   0          30m
sentiment-deployment-7fb479685f-xz5sh  1/1     Running   0          66s
```

Результат describe

kubectl describe deployment sentiment-deployment

kubectl describe svc sentiment-service

kubectl describe ingress sentiment-ingress

kubectl get events --sort-by=.metadata.creationTimestamp

```
Name:                   sentiment-deployment
Namespace:              default
CreationTimestamp:      Mon, 24 Nov 2025 23:05:26 +0300
Labels:                 <none>
Annotations:            deployment.kubernetes.io/revision: 1
Selector:               app=sentiment
Replicas:               6 desired | 6 updated | 6 total | 6 available | 0 unavailable
StrategyType:           RollingUpdate
MinReadySeconds:        0
RollingUpdateStrategy:  25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=sentiment
  Containers:
   sentiment-app:
    Image:      sentiment-app:1.0.0
    Port:       8080/TCP
    Host Port:  0/TCP
    Limits:
      cpu:     500m
      memory:  512Mi
    Requests:
      cpu:         100m
      memory:      256Mi
    Environment:   <none>
    Mounts:        <none>
  Volumes:         <none>
  Node-Selectors:  <none>
  Tolerations:     <none>
Conditions:
  Type           Status  Reason
  ----           ------  ------
```

```
  ----           ------  ------
  Progressing    True    NewReplicaSetAvailable
  Available      True    MinimumReplicasAvailable
OldReplicaSets:  <none>
NewReplicaSet:   sentiment-deployment-7fb479685f (6/6 replicas created)
Events:
  Type    Reason             Age    From                   Message
  ----    ------             ----   ----                   -------
  Normal  ScalingReplicaSet  33m    deployment-controller  Scaled up replica set sentiment-deployment-7fb479685f from 0 to 3
  Normal  ScalingReplicaSet  4m13s  deployment-controller  Scaled up replica set sentiment-deployment-7fb479685f from 3 to 6
Warning: v1 Endpoints is deprecated in v1.33+; use discovery.k8s.io/v1 EndpointSlice
Name:                     sentiment-service
Namespace:                default
Labels:                   <none>
Annotations:              <none>
Selector:                 app=sentiment
Type:                     LoadBalancer
IP Family Policy:         SingleStack
IP Families:              IPv4
IP:                       10.108.94.229
IPs:                      10.108.94.229
Port:                     <unset>  80/TCP
TargetPort:               8080/TCP
NodePort:                 <unset>  32110/TCP
Endpoints:                10.244.0.13:8080,10.244.0.14:8080,10.244.0.15:8080 + 3 more...
Session Affinity:         None
External Traffic Policy:  Cluster
Events:                   <none>
Warning: v1 Endpoints is deprecated in v1.33+; use discovery.k8s.io/v1 EndpointSlice
Name:             sentiment-ingress
Labels:           <none>
Namespace:        default
Address:          192.168.49.2
Ingress Class:    nginx
```

```
  Host            Path  Backends
  ----            ----  --------
  sentiment.local
                  /api   sentiment-service:80 (10.244.0.13:8080,10.244.0.14:8080,10.244.0.15:8080 + 3 more...)
Annotations:      <none>
Events:
  Type    Reason  Age               From                      Message
  ----    ------  ----              ----                      -------
  Normal  Sync    32m (x2 over 33m)  nginx-ingress-controller  Scheduled for sync
LAST SEEN   TYPE      REASON                       OBJECT                                    MESSAGE
33m         Normal    Scheduled                    pod/sentiment-deployment-7fb479685f-mrnss    Successfully assigned default/sentiment-deployment-7fb479685f-mrnss to minikube
33m         Normal    ScalingReplicaSet            deployment/sentiment-deployment           Scaled up replica set sentiment-deployment-7fb479685f from 0 to 3
33m         Normal    SuccessfulCreate             replicaset/sentiment-deployment-7fb479685f   Created pod: sentiment-deployment-7fb479685f-mrnss
33m         Normal    SuccessfulCreate             replicaset/sentiment-deployment-7fb479685f   Created pod: sentiment-deployment-7fb479685f-xsphc
33m         Normal    SuccessfulCreate             replicaset/sentiment-deployment-7fb479685f   Created pod: sentiment-deployment-7fb479685f-w96zm
33m         Normal    Scheduled                    pod/sentiment-deployment-7fb479685f-xsphc    Successfully assigned default/sentiment-deployment-7fb479685f-xsphc to minikube
33m         Normal    Scheduled                    pod/sentiment-deployment-7fb479685f-w96zm    Successfully assigned default/sentiment-deployment-7fb479685f-w96zm to minikube
32m         Normal    Sync                         ingress/sentiment-ingress                 Scheduled for sync
33m         Normal    Created                      pod/sentiment-deployment-7fb479685f-mrnss    Created container: sentiment-app
33m         Normal    Created                      pod/sentiment-deployment-7fb479685f-w96zm    Created container: sentiment-app
33m         Normal    Started                      pod/sentiment-deployment-7fb479685f-xsphc    Started container sentiment-app
33m         Normal    Created                      pod/sentiment-deployment-7fb479685f-xsphc    Created container: sentiment-app
33m         Normal    Pulled                       pod/sentiment-deployment-7fb479685f-xsphc    Container image "sentiment-app:1.0.0" already present on machine
33m         Normal    Started                      pod/sentiment-deployment-7fb479685f-mrnss    Started container sentiment-app
33m         Normal    Pulled                       pod/sentiment-deployment-7fb479685f-mrnss    Container image "sentiment-app:1.0.0" already present on machine
33m         Normal    Started                      pod/sentiment-deployment-7fb479685f-w96zm    Started container sentiment-app
33m         Normal    Pulled                       pod/sentiment-deployment-7fb479685f-w96zm    Container image "sentiment-app:1.0.0" already present on machine
32m         Warning   FailedGetResourceMetric      horizontalpodautoscaler/sentiment-hpa     failed to get cpu utilization: unable to get metrics for resource cpu: no metrics returned
 from resource metrics API
32m         Warning   FailedComputeMetricsReplicas   horizontalpodautoscaler/sentiment-hpa     invalid metrics (1 invalid out of 1), first error is: failed to get cpu resource metric va
lue: failed to get cpu utilization: unable to get metrics for resource cpu: no metrics returned from resource metrics API
31m         Warning   FailedGetResourceMetric      horizontalpodautoscaler/sentiment-hpa     failed to get cpu utilization: did not receive metrics for targeted pods (pods might be un
ready)
31m         Warning   FailedComputeMetricsReplicas   horizontalpodautoscaler/sentiment-hpa     invalid metrics (1 invalid out of 1), first error is: failed to get cpu resource metric va
```

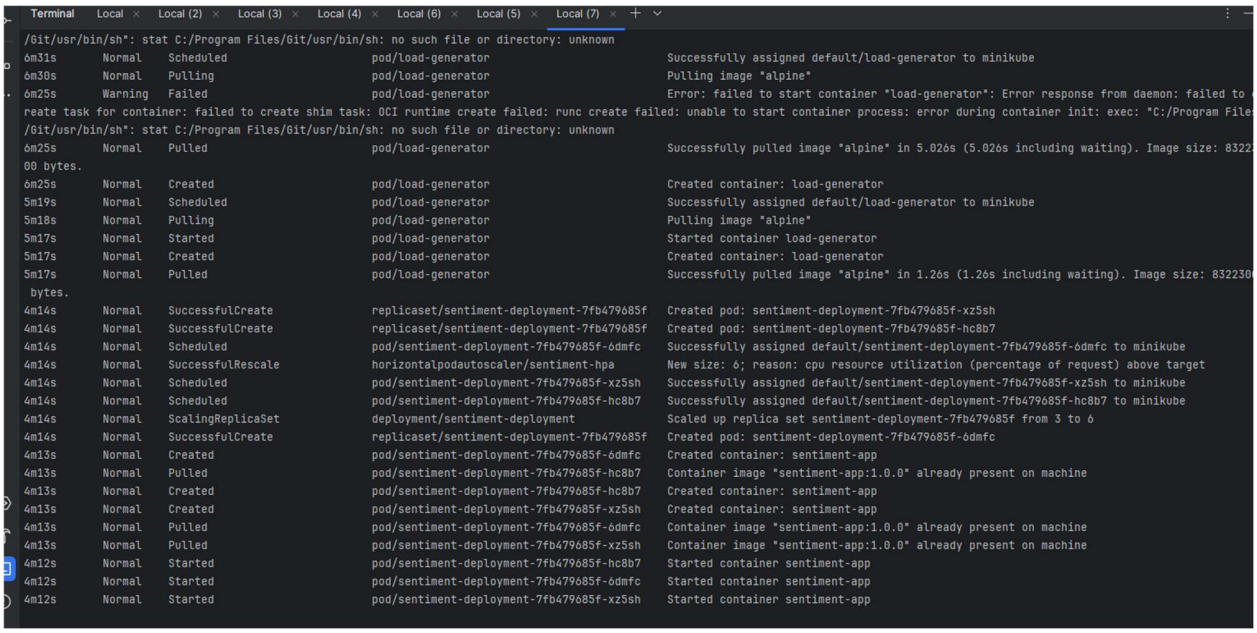
Terminal
Local
Local (2)
Local (3)
Local (4)
Local (6)
Local (5)
Local (7)
/Git/usr/bin/sh*: stat C:/Program Files/Git/usr/bin/sh: no such file or directory: unknown
6m31s Normal Scheduled pod/load-generator Successfully assigned default/load-generator to minikube
6m30s Normal Pulling pod/load-generator Pulling image "alpine"
6m25s Warning Failed pod/load-generator Error: failed to start container "load-generator": Error response from daemon: failed to
reate task for container: failed to create shim task: OCI runtime create failed: runc create failed: unable to start container process: error during container init: exec: "C:/Program File
/Git/usr/bin/sh*: stat C:/Program Files/Git/usr/bin/sh: no such file or directory: unknown
6m25s Normal Pulled pod/load-generator Successfully pulled image "alpine" in 5.026s (5.026s including waiting). Image size: 8322
00 bytes.
6m25s Normal Created pod/load-generator Created container: load-generator
5m19s Normal Scheduled pod/load-generator Successfully assigned default/load-generator to minikube
5m18s Normal Pulling pod/load-generator Pulling image "alpine"
5m17s Normal Started pod/load-generator Started container load-generator
5m17s Normal Created pod/load-generator Created container: load-generator
5m17s Normal Pulled pod/load-generator Successfully pulled image "alpine" in 1.26s (1.26s including waiting). Image size: 832230
bytes.
4m14s Normal SuccessfulCreate replicaset/sentiment-deployment-7fb479685f Created pod: sentiment-deployment-7fb479685f-xz5sh
4m14s Normal SuccessfulCreate replicaset/sentiment-deployment-7fb479685f Created pod: sentiment-deployment-7fb479685f-hc8b7
4m14s Normal Scheduled pod/sentiment-deployment-7fb479685f-ódmfc Successfully assigned default/sentiment-deployment-7fb479685f-ódmfc to minikube
4m14s Normal SuccessfulRescale horizontalpodautoscaler/sentiment-hpa New size: ó; reason: cpu resource utilization (percentage of request) above target
4m14s Normal Scheduled pod/sentiment-deployment-7fb479685f-xz5sh Successfully assigned default/sentiment-deployment-7fb479685f-xz5sh to minikube
4m14s Normal Scheduled pod/sentiment-deployment-7fb479685f-hc8b7 Successfully assigned default/sentiment-deployment-7fb479685f-hc8b7 to minikube
4m14s Normal ScalingReplicaSet deployment/sentiment-deployment Scaled up replica set sentiment-deployment-7fb479685f from 3 to ó
4m14s Normal SuccessfulCreate replicaset/sentiment-deployment-7fb479685f Created pod: sentiment-deployment-7fb479685f-ódmfc
4m13s Normal Created pod/sentiment-deployment-7fb479685f-ódmfc Created container: sentiment-app
4m13s Normal Pulled pod/sentiment-deployment-7fb479685f-hc8b7 Container image "sentiment-app:1.0.0" already present on machine
4m13s Normal Created pod/sentiment-deployment-7fb479685f-hc8b7 Created container: sentiment-app
4m13s Normal Created pod/sentiment-deployment-7fb479685f-xz5sh Created container: sentiment-app
4m13s Normal Pulled pod/sentiment-deployment-7fb479685f-ódmfc Container image "sentiment-app:1.0.0" already present on machine
4m13s Normal Pulled pod/sentiment-deployment-7fb479685f-xz5sh Container image "sentiment-app:1.0.0" already present on machine
4m12s Normal Started pod/sentiment-deployment-7fb479685f-hc8b7 Started container sentiment-app
4m12s Normal Started pod/sentiment-deployment-7fb479685f-ódmfc Started container sentiment-app
4m12s Normal Started pod/sentiment-deployment-7fb479685f-xz5sh Started container sentiment-app

10. Мониторинг — minikube dashboard

minikube dashboard

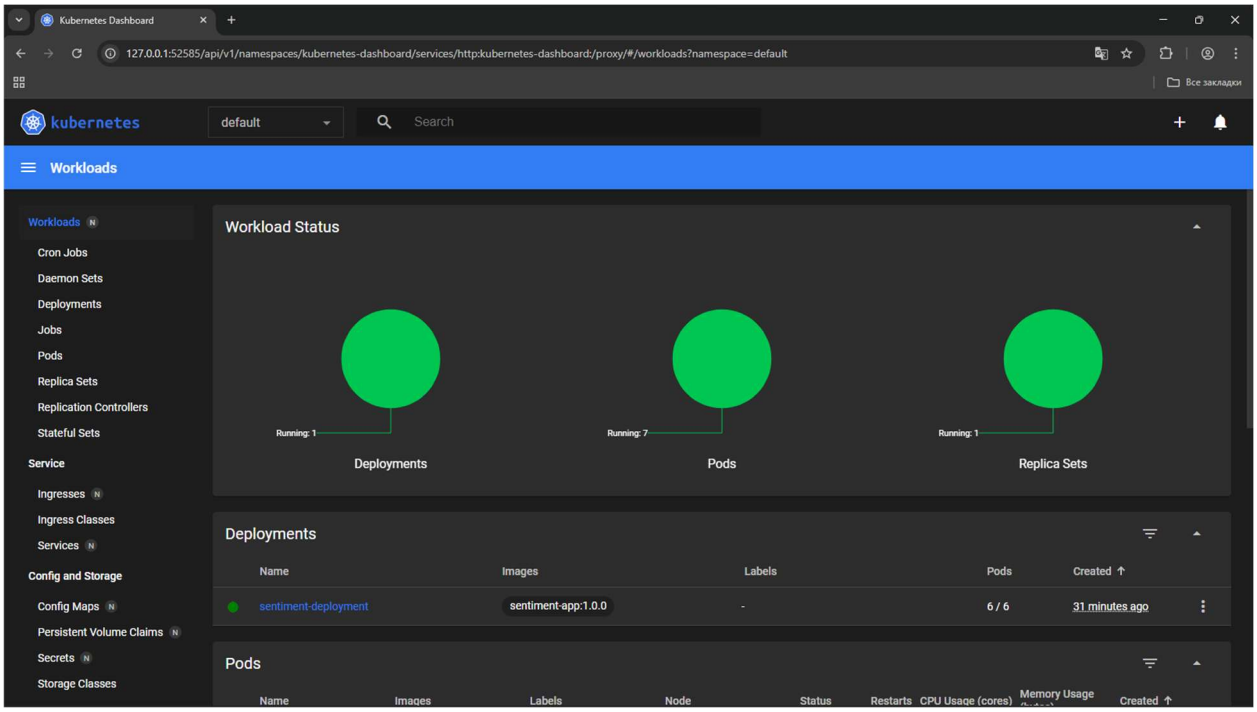

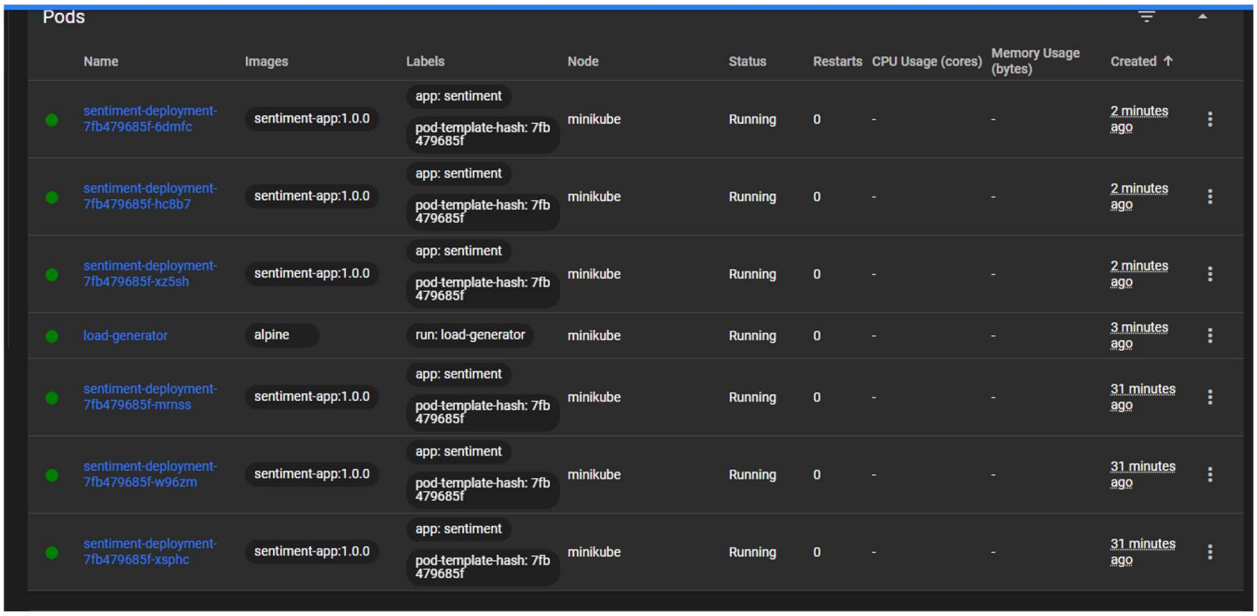

12. Заключение

В ходе работы был полностью реализован проект по разработке и развёртыванию ИИ-приложения в Kubernetes:

1. создано Spring Boot приложение с REST API и ИИ-логикой анализа тональности;
2. приложение успешно контейнеризировано (образ <150 МБ);
3. развернуто в Minikube с использованием Deployment + Service + Ingress;
4. добавлены monitoring endpoints через Actuator;
5. настроен HPA и проведён стресс-тест;
6. добились масштабирования с 3 до 6 реплик при нагрузке 194%;